

CS 201

Computer Systems Programming

Instructor: A.Prof. Dinh Dien

This lecture is modified from course CS201 of Portland State University

Hello, World

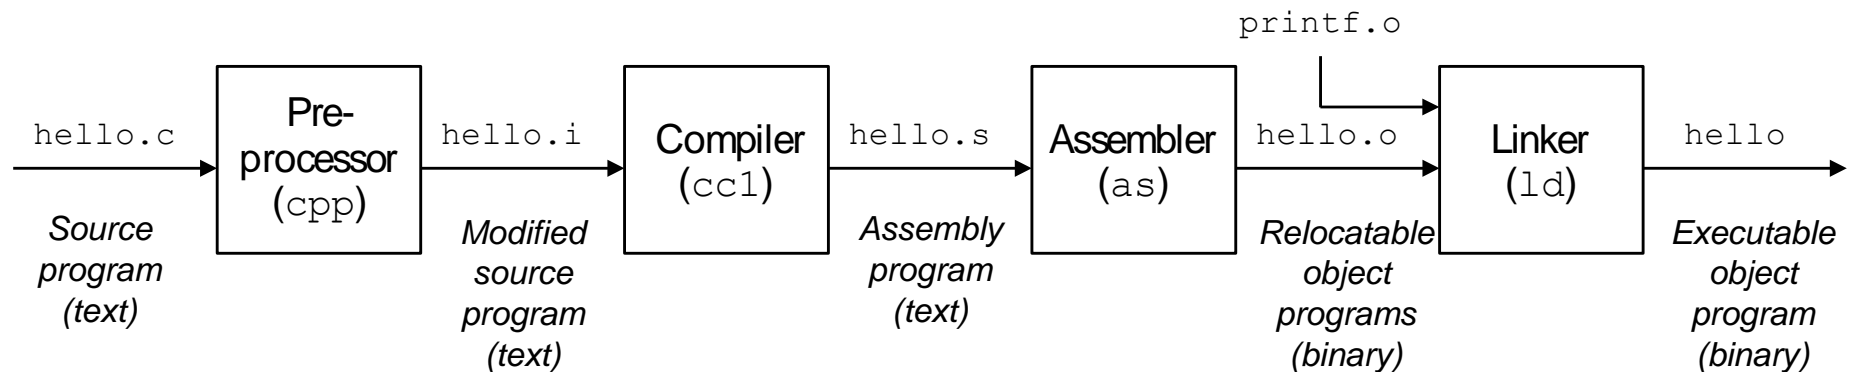
```
#include <stdio.h>
```

```
int main()
```

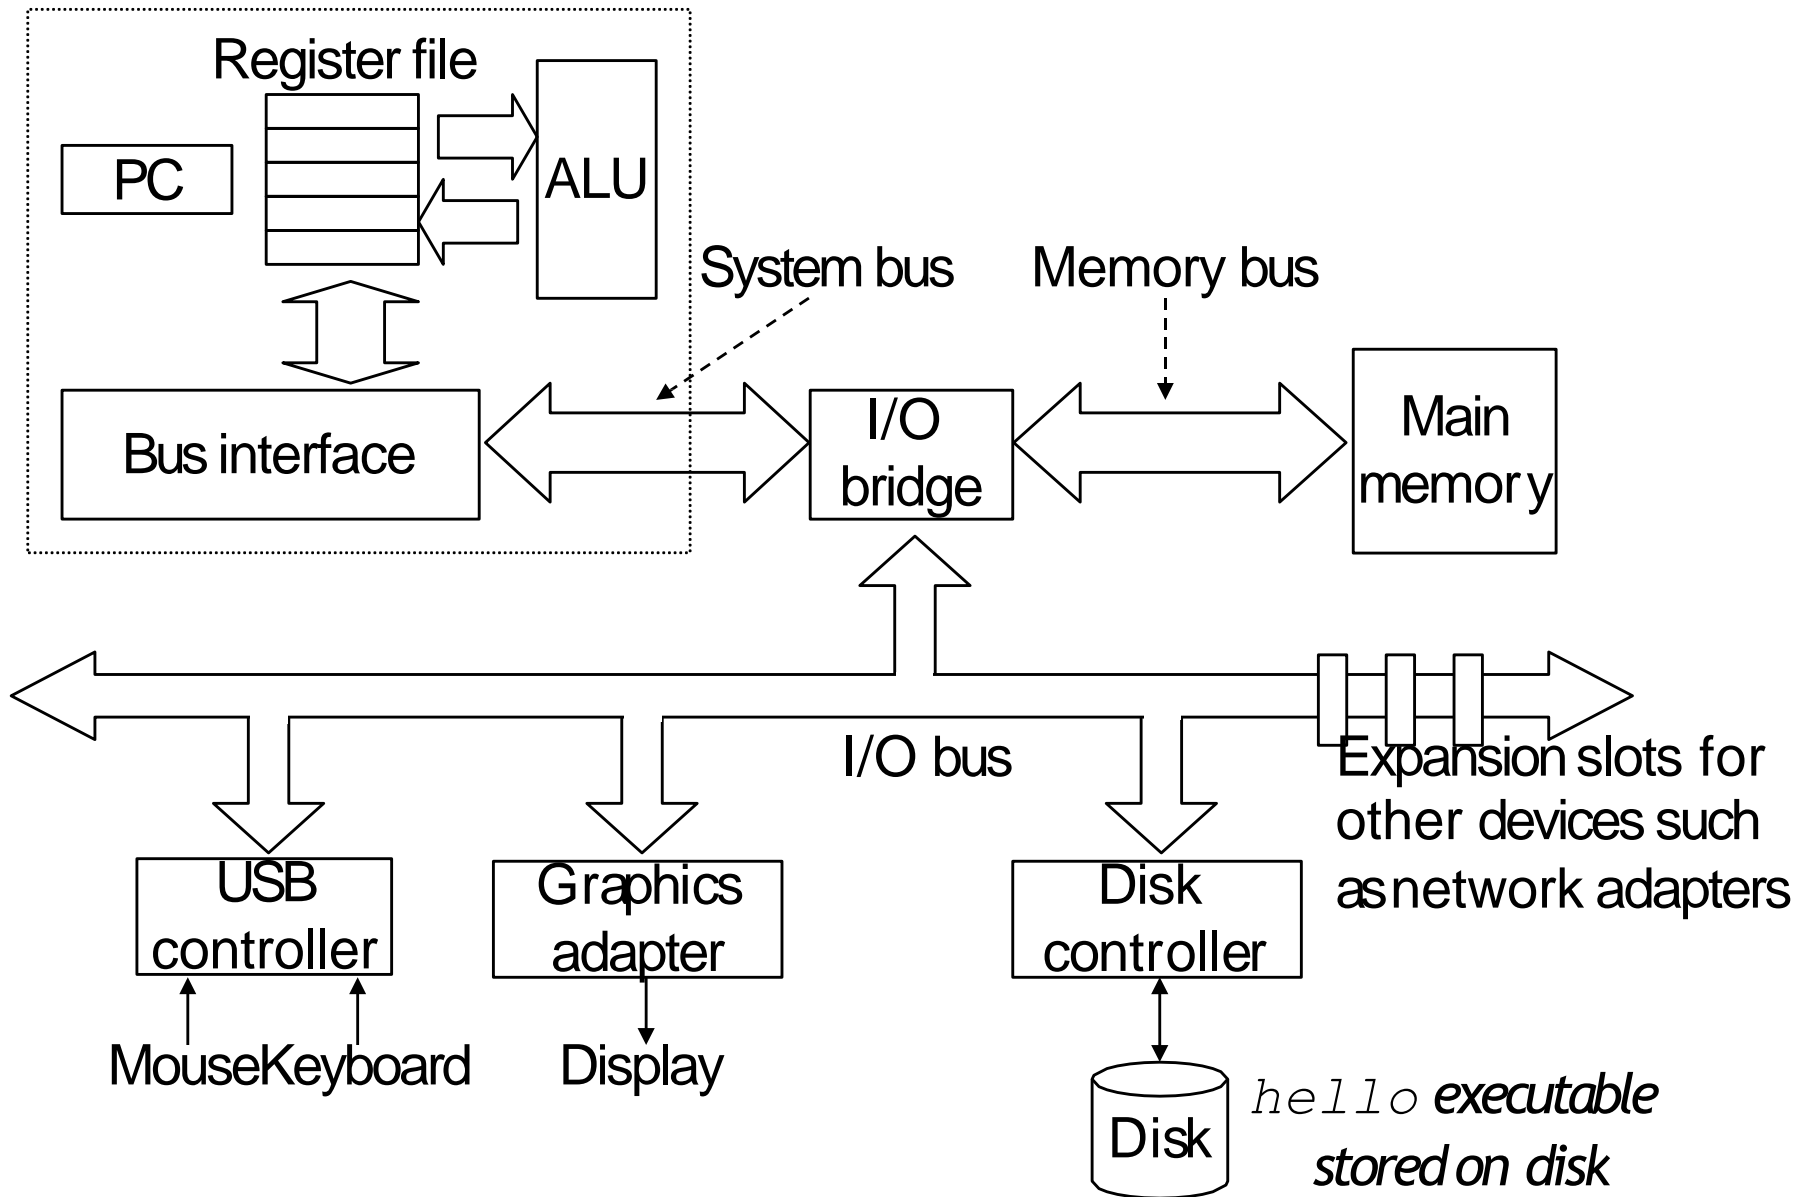
```
{
```

```
    printf("hello, world\n");
```

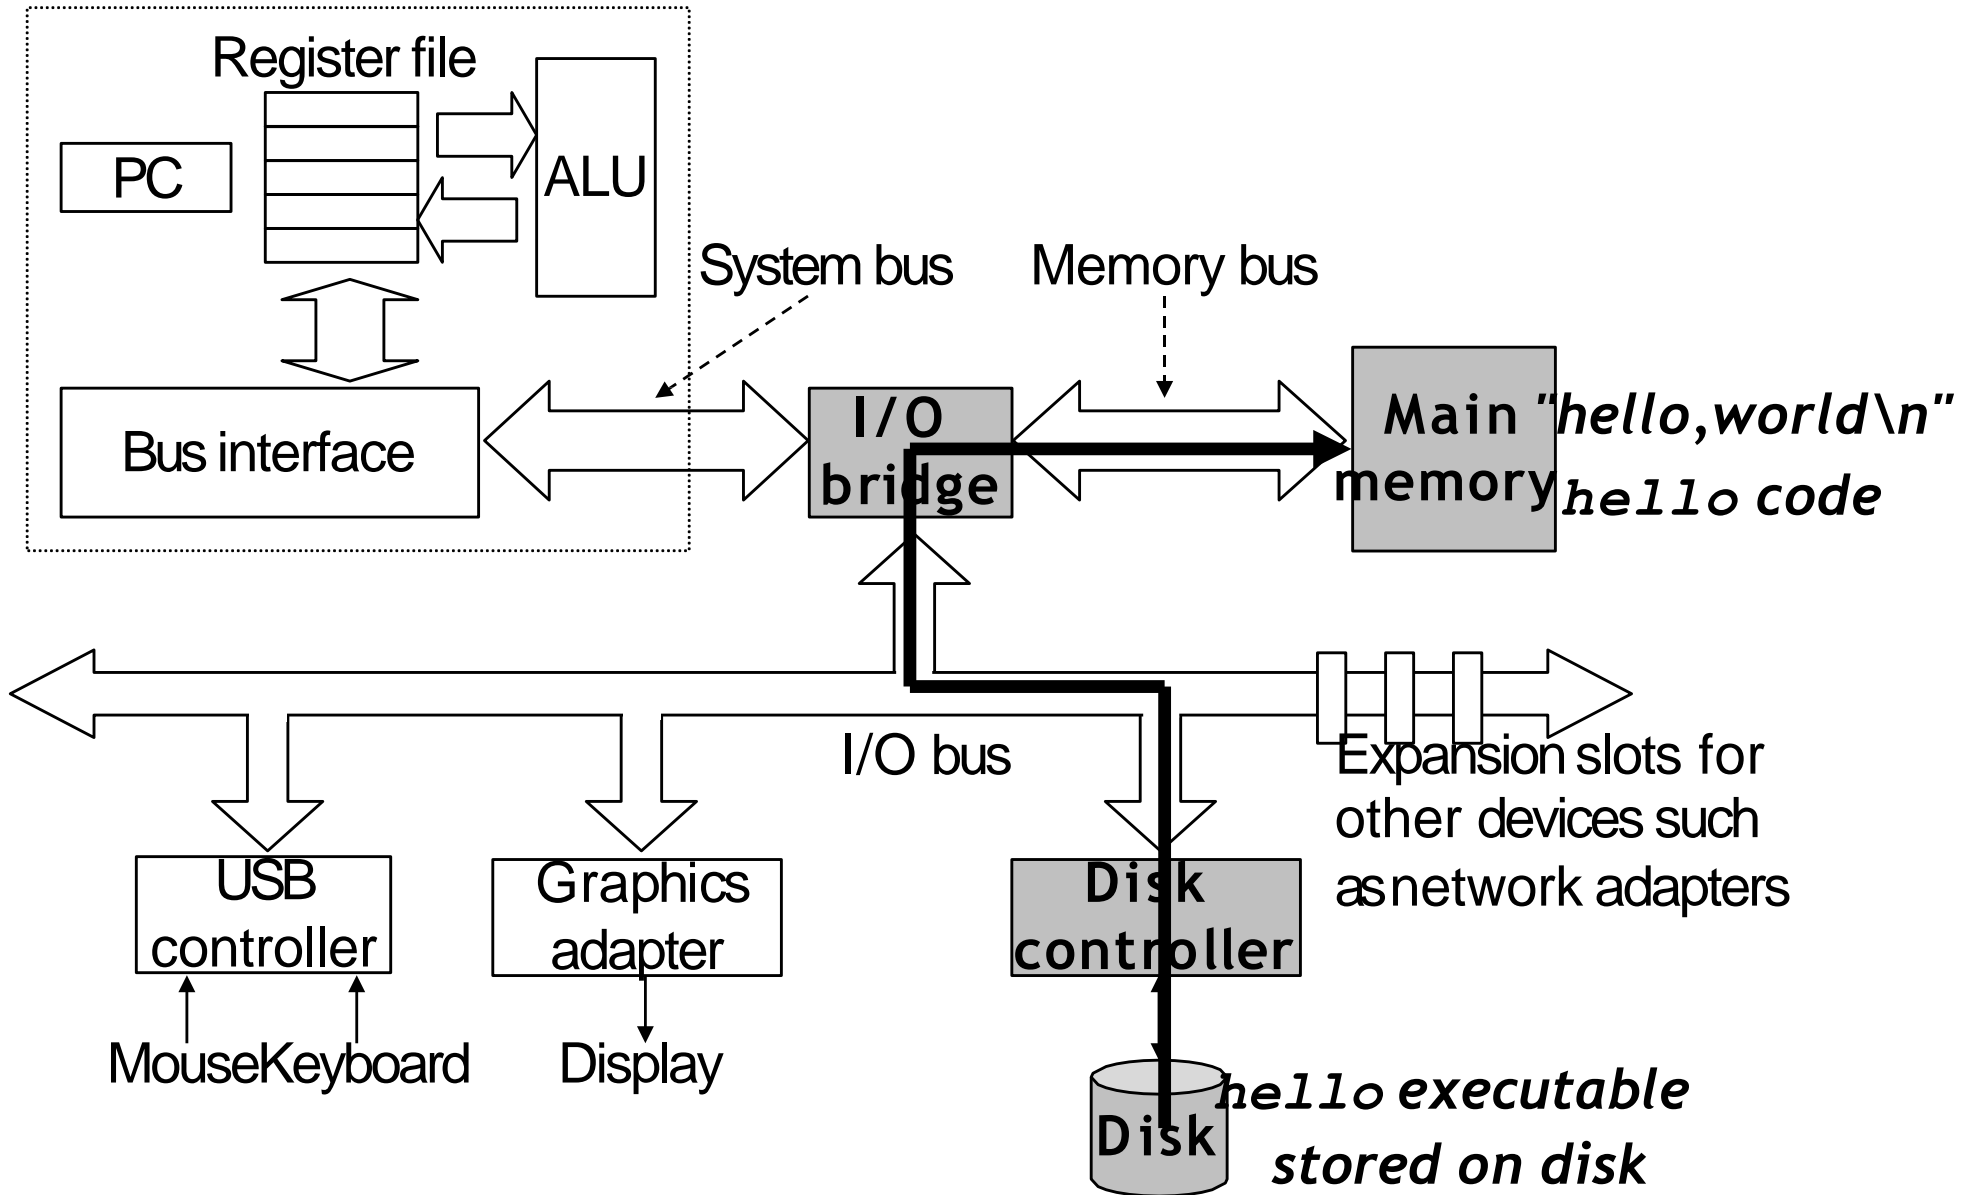
```
}
```



CPU



CPU



Course Theme:

Abstraction Is Good But Don't Forget Reality

- ❑ **Most CS and CE courses emphasize abstraction**
 - Abstract data types
 - Processes, Files
- ❑ **These abstractions have limits**
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations
- ❑ **Useful outcomes**
 - Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
 - Prepare for later “systems” classes in CS & ECE
 - Compilers, Operating Systems, Networks, Architecture

Great Reality #1:


Ints are not Integers, Floats are not Reals

Example 1: Is $x^2 \geq 0$?

Float's: Yes!

Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ??$



Example 2: Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!
- Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Code Security Example

```
/* Kernel memory region holding user-accessible data
*/  #define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer
*/  int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen
    */  int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

- ⌘ **Similar to code found in FreeBSD's implementation of `getpeername`**
- ⌘ **There are legions of smart people trying to find vulnerabilities in programs**

Typical Usage

```
/* Kernel memory region holding user-accessible data
*/  #define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer
*/  int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen
    */  int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```


Malicious Usage

```
/* Kernel memory region holding user-accessible data
*/  #define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer
*/  int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen
    */  int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    . . .
}
```

Computer Arithmetic

- ❑ **Does not generate random values**
 - Arithmetic operations have important mathematical properties
- ❑ **Cannot assume all “usual” mathematical properties**
 - Due to finiteness of representations
 - Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
 - Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs
- ❑ **Observation**
 - Need to understand which abstractions apply in which contexts
 - Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- ❑ **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- ❑ **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Assembly Code Example

❑ Time Stamp Counter

- Special 64-bit register in Intel-compatible machines
- Incremented every clock cycle
- Read with rdtsc instruction

❑ Application

- Measure time (in clock cycles) required by procedure

```
double t;  
start_counter()  
; P();  
t = get_counter();  
printf("P required %f clock cycles\n", t);
```

Code to Read Counter

- ❏ Write small amount of assembly code using GCC's asm facility
- ❑ Inserts assembly code into machine code generated by compiler

```
static unsigned cyc_hi = 0;
static unsigned cyc_lo = 0;

/* Set *hi and *lo to the high and low order
   bits of the cycle counter.
*/
void access_counter(unsigned *hi, unsigned *lo)
{
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"
        : "=r" (*hi), "=r" (*lo)
        :
        : "%edx", "%eax");
}
```

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

- ❑ **Memory is not unbounded**
 - It must be allocated and managed
 - Many applications are memory dominated
- ❑ **Memory referencing bugs especially pernicious**
 - Effects are distant in both time and space
- ❑ **Memory performance is not uniform**
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
double fun(int i)
{
    volatile double d[1] = {3.14};
    volatile long int a[2];
    a[i] = 1073741824; /* Possibly out of bounds
    */ return d[0];
}
```

fun(0) → 3.14
fun(1) → 3.1399998664856
fun(2) → 3.14
fun(3) → 2.00000061035156
fun(4) → 3.14, then segmentation fault

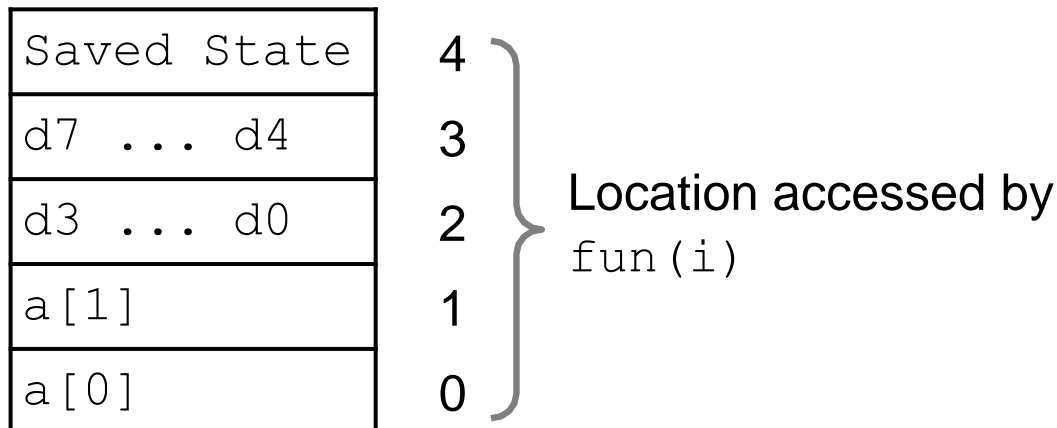
❑ **Result is architecture specific**

Memory Referencing Bug Example

```
double fun(int i)
{
    volatile double d[1] = {3.14};
    volatile long int a[2];
    a[i] = 1073741824; /* Possibly out of bounds
    */ return d[0];
}
```

fun(0) → 3.14
 fun(1) → 3.14
 fun(2) → 3.13999998664856
 fun(3) → 2.000000061035156
 fun(4) → 3.14, then segmentation fault

Explanation:




Memory Referencing Errors

- ❑ **C and C++ do not provide any memory protection**
 - Out of bounds array references
 - Invalid pointer values
 - Abuses of malloc/free
- ❑ **Can lead to nasty bugs**
 - Whether or not bug has any effect depends on system and compiler
 - Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated
- ❑ **How can I deal with this?**
 - Program in Java, Ruby or ML
 - Understand what possible interactions may occur
 - Use or develop tools to detect referencing errors (e.g. Valgrind)

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```




21 times slower
(Pentium 4)

- ❑ Hierarchical memory organization
- ❑ Performance depends on access patterns
 - Including how step through multi-dimensional array

The Memory Mountain

Intel Core i7
2.67 GHz
32 KB L1 d-cache
256 KB L2 cache
8 MB L3 cache




Great Reality #4: There's more to performance than asymptotic complexity

- ❑ **Constant factors matter too!**
- ❑ **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- ❑ **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Example Matrix Multiplication

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz (double precision)
Gflop/s




- ❑ Standard desktop computer, vendor compiler, using optimization flags
- ❑ Both implementations have **exactly** the same operations count ($2n^3$)
- ❑ **What is going on?**

MMM Plot: Analysis

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz

Gflop/s



- Reason for 20x: Blocking or tiling, loop unrolling, array scalarization, instruction scheduling, search to find best choice
- Effect: fewer register spills, L1/L2 cache misses, and TLB misses**

Great Reality #5:

Computers do more than execute programs

- ❑ **They need to get data in and out**
 - I/O system critical to program reliability and performance
 - Data movement is currently seen as our biggest obstacle to pushing the high end of computing even higher

- ❑ **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Great Reality #6: Security Matters!!!

- ❑ If you do not understand this level of programming, you are at risk to introduce security holes in code you write

Required Course Textbook

- ❑ **Randal E. Bryant and David R. O'Hallaron,**
 - “Computer Systems: A Programmer’s Perspective, Third Edition” (CS:APP3e), Prentice Hall, 2015
 - Additional student materials at: <http://csapp.cs.cmu.edu>
 - NOTE: BRAND NEW BOOK
 - Key change from 2nd edition: focus on 64-bit architecture
 - “what’s 64-bit architecture?” - stay tuned!
- ❑ **C Programming notes available online for free:**
<https://www.eskimo.com/~scs/cclass/cclass.html>

Getting Help

- ❑ **Class Web Page:** <http://web.cecs.pdx.edu/~karavan/cs201>
- ❑ **We will transition to D2L**
- ❑ **TA Office Hours:** **TBD**
- ❑ **Instructor office hour:** **TBD**
- ❑ **Appointments are always possible modulo my schedule:**
 - Email me and list some different days/times you can meet
 - If you cannot make it – please email and let me know if at all possible
- ❑ **“Open Door” Policy**
 - I am in my office and the door is open == Welcome !
 - I am in my office and the door is [latched] shut == Oops! Not Now.
 - I am in a meeting or on a deadline and request no interruptions other than life and death emergencies

Facilities

❑ Labs and homeworks will use the Computer Science Linux Lab

- ❑ Remote login: [ssh myloginname@linuxlab.cs.pdx.edu](ssh://myloginname@linuxlab.cs.pdx.edu)
- CS tutors sit outside of the lab during posted hours
- Small library of relevant books maintained by tutors

❑ Homeworks

- Homeworks will be tested and graded on the Linux Lab machines
- We do not have the resources to accommodate your individual personal machine setups – please test your work on the lab machines before submitting

Timeliness

❑ Grace days

- 2 “free passes” (48 hours each) for the homeworks
- Covers scheduling crunch, out-of-town trips, illnesses, minor setbacks
- Murphy Says: Save them until late in the term!

❑ Lateness penalties

- Once free passes(s) used up, receive a score of 0
- TURN IN WHATEVER YOU HAVE !

❑ Free Lunch

- Your homework score is computed using the BEST 7 of 8 homeworks

❑ Advice

- Once you start running late, it's really hard to catch up
- 8 Weeks goes by VERYVERY quickly

Cheating

❑ What is cheating?

- Sharing code: by copying, retyping, looking at, or supplying a file
- Coaching: helping your friend to write a lab, line by line
- Copying code from previous course or from elsewhere on WWW
 - Only allowed to use code we supply
- Looking at anyone else's exam or showing anyone yours, in the exam room
- Posting in any form or forum the homework or exam answers

❑ What is NOT cheating?

- Explaining how to use systems or tools or getting that explained
- Helping others with high-level design issues or getting that help
- Getting help from the tutors or Course Expert

❑ Murphy Says:

- Tends to happen when you're tired, behind, and worried – so stay on track
- 2 min rule: after an explanation, 2mins before hands are back on keyboard
- “Why are we spending all this time talking about cheating? “
 - Because it happens EVERY YEAR and if it happens to you, you will FAIL

Other Rules of the CS 201 Classroom

- ❑ Laptops: permitted
- ❑ Electronic communications: *forbidden*
 - No email, instant messaging, cell phone calls, etc
 - You will be asked to leave
 - Why? This is *a Learning Environment*
 - OK- *looking things up as we go, electronic textbook, etc.*
 - OK – *asking questions in the class*
- ❑ ***Note: Some students have been granted specific permission to record the lectures. Without permission it's a no-no.***

HW #1 Part 1

- ❑ The Full Homework will be available on Thursday (April 4)
- ❑ But it's a good idea to get started

1. Readings:

B&O [textbook] chapter 1

Steve Summit's C Programming Notes Chapter 1:

<https://www.eskimo.com/~scs/cclass/notes/top.html>

1) Hands on:

- Online Tutorial (SKIP 105.7): <http://pages.cs.wisc.edu/~remzi/OSTEP/lab-tutorial.pdf>
 - Focus on your initial goals: to be able to write, compile, run C Programs on the PSU Linux Lab machines
- *The Key*: Learn a command line editor: vim or emacs

*Welcome
and Enjoy!*